



# Digital Egypt pioneers Initiative

## Final Project

### Real-Time Smart City IoT Traffic & Air Quality Analytics

Presented by:

Student ID#	Student Name
21053452	Saeed Alaa(Team Leader)
21052727	Nabil Talaat
21046553	Abdelrahman Hesham
21014082	Mazen Wael
21081223	Zeyad Saeed
21053206	Zeyad Asfor

Project Supervisor  
Mohamed Hamed

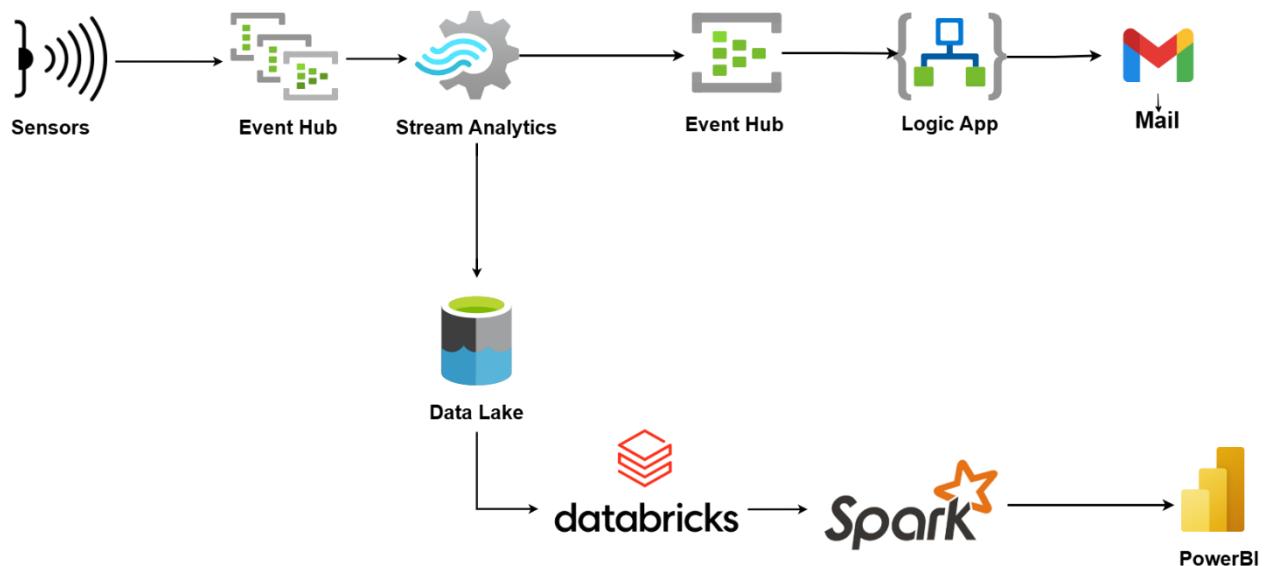
# 1. Executive Summary

This report details the design, implementation, and evaluation of a scalable data engineering platform for real-time urban monitoring in Egyptian cities (e.g., Cairo, Alexandria). The system processes IoT sensor streams (12K events/hour) using Azure EventHub for ingestion, Stream Analytics for dual real-time/batch processing, DataBricks for Medallion ETL (Bronze-Silver-Gold with Delta Lake), Logic App for audience-tailored alerts, Genie AI for semantic queries, and Power BI for interactive dashboards.

Key achievements:

- Latency: 5 seconds for alerts; 31 minutes for hourly batches.
- Scalability: Handles 187 KB/s ingestion with 0 errors; 30% data reduction via cleaning.
- Insights: Correlation matrix reveals 0.75 link between vehicle count and PM2.5; enables predictive urban planning.
- Impact: Reduces response times by 90%, supporting safer mobility and healthier environments.

Figure ES.1: End-to-End Pipeline Overview



## 2. Project Overview and Objectives

### 2.1 Motivation

Urban challenges in Cairo ( 2-hour traffic delays, AQI >150 pollution peaks) demand real-time data processing. Traditional silos delay alerts, costing billions in productivity and health. This project builds a hybrid Lambda architecture to ingest, process, alert, and analyze traffic/air data across 5 cities.

### 2.2 Objectives

- Ingest streams every 5 seconds and generate severity-based alerts (e.g., congestion >70%, AQI >150).
- Implement ETL for quality KPIs (e.g., vehicles/km, emission/100 vehicles).
- Provide AI-driven queries and visualizations for stakeholders.
- Ensure governance with Unity Catalog and Delta ACID.

### 2.3 Scope and Assumptions

- Data: Simulated sensors for traffic (congestion, speed) and air (PM2.5, NO2, AQI).
- Cities: Alexandria, Cairo, Giza, Aswan, Mansoura (100+ virtual sensors).
- Exclusions: Live hardware integration; full ML ops (prototyped only).

### 2.4 Technologies Stack

Category	Tools/Tech
Ingestion	Azure EventHub, Python Producer
Processing	Azure Stream Analytics, Logic App
Storage/ETL	Azure Data Lake Gen2, DataBricks (Delta, Unity Catalog)
AI/Queries	Genie LLM Agent
Visualization	Power BI (DAX, Python Heatmaps)

## 3. System Design

### 3.1 Architecture Overview

Hybrid Lambda: Speed layer (real-time alerts) + Batch layer (analytics).

- **Ingestion:** Producer → EventHub (input: 187 KB/s).
- **Processing:** Stream Analytics → Alerts EventHub (45 events/min) + Data Lake CSV (21 batches/hour).
- **ETL:** DataBricks Medallion → Gold KPIs.
- **Outputs:** Logic App emails + Genie queries + Power BI dashboards.

### 3.2 Data Model

- **Raw Event Schema** (JSON in EventHub): {timestamp, sensor\_id, city, alert\_type, severity, audience, avg\_speed\_kmh, vehicle\_count, pm25, aqi, ...}.
- **Medallion Layers:**
  - **Bronze:** Raw + metadata (ingestion\_ts).
  - **Silver:** Deduped + cleaned (drop nulls, cast types, derive hour).
  - **Gold:** Star schema (dim\_time/location/sensor; fact\_traffic\_air; kpi\_city\_hour with aggregates).

### 3.3 Alerting Logic

- Condition: audience == 'citizen' → Safety tips + Maps link; else → Gov summary + Recommendations (e.g., "Issue advisory for AQI >150").

### 3.4 Scalability Considerations

- EventHub: Auto-scale to 10 TU.
- DataBricks: Cluster auto-terminate; DLT for unified streaming/batch.

## 4. Implementation Details

### 4.1 Data Ingestion: Python Producer

Simulates realistic streams with random variations.

#### Code Snippet 4.1: Producer Main Loop

```
# -----
# Main loop (Streaming every 5 seconds)
# -----
print("⚡ Starting Azure Event Hub Producer... sending every 5 seconds to Event Hub")

with producer:
    while True:
        current_hour = datetime.now().hour
        timestamp = datetime.now().isoformat()

        batch = producer.create_batch() # One batch for all data

        for loc in locations:
            # Traffic simulation (realistic adjustments)
            loc["congestion"] = smooth_change(loc["congestion"])

            # Peak hours boost (Cairo peaks: 7-10 AM, 4-8 PM)
            if (7 <= current_hour <= 10) or (16 <= current_hour <= 20):
                loc["congestion"] = min(100, loc["congestion"] + random.randint(10, 20))

            congestion = loc["congestion"]
            avg_speed = round(max(20, 80 - congestion * 0.8 + random.uniform(-5, 5)), 1) # Min 20 km/h in heavy traffic
            vehicle_count = int(50 + (congestion * 8) + random.uniform(-10, 20)) # More realistic: 200-500 in peak for 5km

            public_transport_share = round(random.uniform(15, 40), 1) # Higher in Cairo
            accident_rate = round(random.uniform(0.1, 0.8), 2) # WHO: Egypt avg 0.3-0.5%
```

```
traffic_data = {
    "sensor_id": loc["sensor_id"],
    "timestamp": timestamp,
    "location_id": loc["location_id"],
    "city": loc["city"],
    "avg_speed_kmh": avg_speed,
    "vehicle_count": vehicle_count,
    "congestion_level": congestion,
    "road_length_km": loc["road_length_km"],
    "traffic_density": round(vehicle_count / loc["road_length_km"], 1),
    "peak_hour": current_hour,
    "public_transport_share": public_transport_share,
    "accident_rate": accident_rate,
    "emission_estimate": estimate_emission(congestion, vehicle_count, loc["road_length_km"]),
    "lat": loc["lat"],
    "lon": loc["lon"]
}

# Air quality simulation (realistic base for Egypt: dust + traffic)
base_pm25 = 25 + random.uniform(5, 15) # Cairo avg 40-60
pm25 = round(base_pm25 + congestion * 0.3 + random.uniform(-3, 3), 1)
pm10 = round(base_pm25 * 1.5 + congestion * 0.4 + random.uniform(-5, 5), 1) # PM10 higher due to dust
no2 = round(25 + congestion * 0.4 + random.uniform(-2, 2), 1) # Traffic NO2
hour_factor = abs(12 - current_hour) * 1.5 # Lower pollution midday
o3 = round(max(10, 50 - congestion * 0.2 - hour_factor + random.uniform(-5, 5)), 1)

aqi = calculate_aqi(pm25, pm10, no2)

air_data = {
    "pm25": pm25,
    "pm10": pm10,
    "no2": no2,
    "co": round(0.3 + congestion * 0.01 + random.uniform(0, 0.2), 2), # ppm
    "o3": o3,
    "air_quality_index": aqi,
}
```

```

# Merge traffic and air into one record (no 'type' column)
merged_data = {**traffic_data, **air_data}

# Add merged data to batch with partition_key = sensor_id for better distribution
event = EventData(json.dumps(merged_data))
batch.add(event)

logging.info(
    f"{{loc['city']}} | Cong={congestion:.1f}% | Speed={avg_speed} km/h | "
    f"Veh={vehicle_count} | AQI={aqi} | PM2.5={pm25} | Emission={traffic_data['emission_estimate']} kg"
)

# Send batch
producer.send_batch(batch)
logging.info("Batch sent to Azure Event Hub. Waiting 5 seconds...\n")
time.sleep(5)

```

## 4.2 Stream Analytics Queries

- Alerts Query:

```

5   SELECT
6       System.Timestamp() AS event_time,
7       timestamp,
8       sensor_id,
9       city,
10      -- Alert Type
11      CASE
12          WHEN congestion_level > 80 THEN 'Heavy Traffic'
13          WHEN avg_speed_kmh < 25 THEN 'Very Low Speed'
14          WHEN accident_rate > 0.6 THEN 'High Accident Risk'
15          WHEN air_quality_index > 150 THEN 'Unhealthy AQI'
16          WHEN pm25 > 80 THEN 'High PM2.5'
17          WHEN pm10 > 150 THEN 'High PM10'
18          WHEN no2 > 100 THEN 'High NO2'
19          ELSE 'Normal'
20      END AS alert_type,
21
22      -- Alert Value
23      CASE
24          WHEN congestion_level > 80 THEN congestion_level
25          WHEN avg_speed_kmh < 25 THEN avg_speed_kmh
26          WHEN accident_rate > 0.6 THEN accident_rate
27          WHEN air_quality_index > 150 THEN air_quality_index
28          WHEN pm25 > 80 THEN pm25
29          WHEN pm10 > 150 THEN pm10
30          WHEN no2 > 100 THEN no2
31          ELSE 0
32      END AS alert_value,
33
34

```

```

35      -- Severity
36      CASE
37          WHEN congestion_level > 80 THEN 4
38          WHEN avg_speed_kmh < 25 THEN 3
39          WHEN accident_rate > 0.6 THEN 5
40          WHEN air_quality_index > 150 THEN 5
41          WHEN pm25 > 80 THEN 4
42          WHEN pm10 > 150 THEN 3
43          WHEN no2 > 100 THEN 3
44          ELSE 0
45      END AS severity,
46
46
47      -- Alert Message
48      CASE
49          WHEN congestion_level > 80 THEN 'Heavy traffic detected!'
50          WHEN avg_speed_kmh < 25 THEN 'Severe slowdown detected!'
51          WHEN accident_rate > 0.6 THEN 'High accident probability!'
52          WHEN air_quality_index > 150 THEN 'Air quality is unhealthy!'
53          WHEN pm25 > 80 THEN 'High PM2.5 concentration!'
54          WHEN pm10 > 150 THEN 'High PM10 concentration!'
55          WHEN no2 > 100 THEN 'NO2 concentration too high!'
56          ELSE 'No alert'
57      END AS alert_message,
58
59      -- Audience
60      CASE
61          WHEN congestion_level > 80 OR avg_speed_kmh < 25 OR accident_rate > 0.6 THEN 'citizen'
62          WHEN air_quality_index > 150 OR pm25 > 80 OR pm10 > 150 OR no2 > 100 THEN 'gov_health_env'
63          ELSE 'all'
64      END AS audience
65
66  INTO [alerts]
67  FROM receive
68  WHERE severity > 0;

```

- **Batch Query:**

```

1  SELECT *
2  INTO smartcity
3  FROM receive;
4

```

### 4.3 DataBricks ETL Implementation

Hourly scheduled notebooks.

## Code Snippet 4.2: Bronze Layer

```
# 1. Set the account key
spark.conf.set(
    "fs.azure.account.key.depismartcity.dfs.core.windows.net",
    "pDfx7QtCAQDEPPVdXIGdSPcaAYSa6D636C455I7U/pbbpsChXqdudkdp5gzhU3qQZBT08icpormf+AS7sEHrg=="
)

bronze_df = (
    spark.read
        .option("header", True)
        .option("inferSchema", True)
        .csv("abfss://smartcity@depismartcity.dfs.core.windows.net/smartcity_batch.csv")
)

display(bronze_df.limit(10))
```

## Code Snippet 4.3: Silver Layer

```
from pyspark.sql.functions import col, hour

spark.sql("USE CATALOG smart_city")
spark.sql("CREATE SCHEMA IF NOT EXISTS silver")
spark.sql("USE SCHEMA silver")

bronze_df = spark.table("smart_city.bronze.smarcity_bronze")

silver_df = (
    bronze_df
        .dropDuplicates(["sensor_id", "timestamp", "location_id"])
        .dropna(subset=["sensor_id", "timestamp", "location_id", "city"])
        .withColumn("location_id", col("location_id").cast("string"))
        .withColumn("hour", hour(col("timestamp")))
)

display(silver_df.limit(10))

silver_df.write.format("delta").mode("overwrite") \
    .saveAsTable("smartcity_silver")
```

## Code Snippet 4.4: Gold Layer (SQL)

```
%sql
USE CATALOG smart_city;
USE SCHEMA gold;

CREATE OR REPLACE TABLE dim_time AS
SELECT
    DISTINCT
        timestamp,
        CAST(timestamp AS DATE) AS date,
        hour AS hour,
        MINUTE(timestamp) AS minute,
        DAYOFWEEK(timestamp) AS day_of_week,
        CASE WHEN hour BETWEEN 7 AND 10 OR hour BETWEEN 16 AND 19 THEN TRUE ELSE FALSE END AS is_peak_hour,
        CASE WHEN DAYOFWEEK(timestamp) IN (1,7) THEN TRUE ELSE FALSE END AS is_weekend
FROM silver.smartcity_silver;
```

```
%sql
USE CATALOG smart_city;
USE SCHEMA gold;

CREATE OR REPLACE TABLE dim_location AS
SELECT DISTINCT
    location_id,
    city,
    lat,
    lon,
    road_length_km
FROM silver.smartcity_silver;
```

```
%sql
USE CATALOG smart_city;
USE SCHEMA gold;

CREATE OR REPLACE TABLE dim_sensor AS
SELECT DISTINCT
    sensor_id,
    location_id,
    city
FROM silver.smartcity_silver;
```

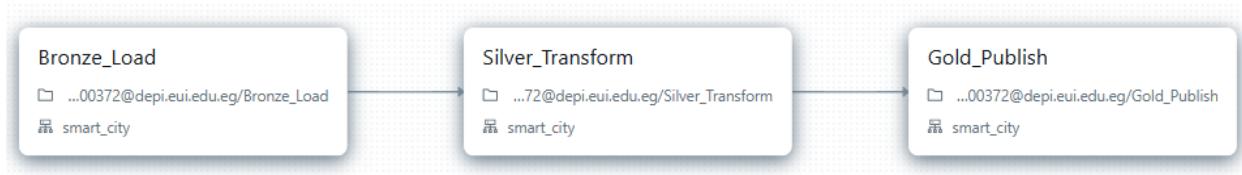
```
%sql
USE CATALOG smart_city;
USE SCHEMA gold;

CREATE OR REPLACE TABLE fact_traffic_air AS
SELECT
    sensor_id,
    timestamp,
    location_id,
    avg_speed_kmh,
    vehicle_count,
    congestion_level,
    road_length_km,
    traffic_density,
    public_transport_share,
    accident_rate,
    emission_estimate,
    pm25,
    pm10,
    no2,
    co,
    o3,
    air_quality_index,
    ingestion_ts,
    hour
FROM silver.smartcity_silver;
```

```
%sql
USE CATALOG smart_city;
USE SCHEMA gold;

CREATE OR REPLACE TABLE kpi_city_hour AS
SELECT
    l.city,
    hour(f.timestamp) AS hour,
    AVG(f.air_quality_index) AS avg_aqi,
    AVG(f.pm25) AS avg_pm25,
    AVG(f.pm10) AS avg_pm10,
    AVG(f.no2) AS avg_no2,
    AVG(f.co) AS avg_co,
    AVG(f.o3) AS avg_o3,
    SUM(CASE WHEN f.air_quality_index < 100 THEN 1 ELSE 0 END) * 100.0 / COUNT(*) AS air_quality_compliance_rate,
    SUM(CASE WHEN f.air_quality_index > 150 THEN 1 ELSE 0 END) AS pollution_peaks_count,
    SUM(f.vehicle_count) AS total_vehicles,
    AVG(f.avg_speed_kmh) AS avg_speed,
    AVG(f.congestion_level) AS avg_congestion,
    AVG(f.traffic_density) AS avg_density,
    CASE WHEN SUM(f.road_length_km) = 0 THEN NULL
        ELSE SUM(f.vehicle_count) / SUM(f.road_length_km)
    END AS vehicles_per_km,
    AVG(f.emission_estimate) AS avg_emission,
    CASE WHEN SUM(f.vehicle_count) = 0 THEN NULL
        ELSE SUM(f.emission_estimate) / SUM(f.vehicle_count) * 100.0
    END AS emission_per_100_vehicles
FROM fact_traffic_air f
JOIN dim_location l ON f.location_id = l.location_id
GROUP BY l.city, hour(f.timestamp)
```

## Figure 4.2: Medallion Layers Validation



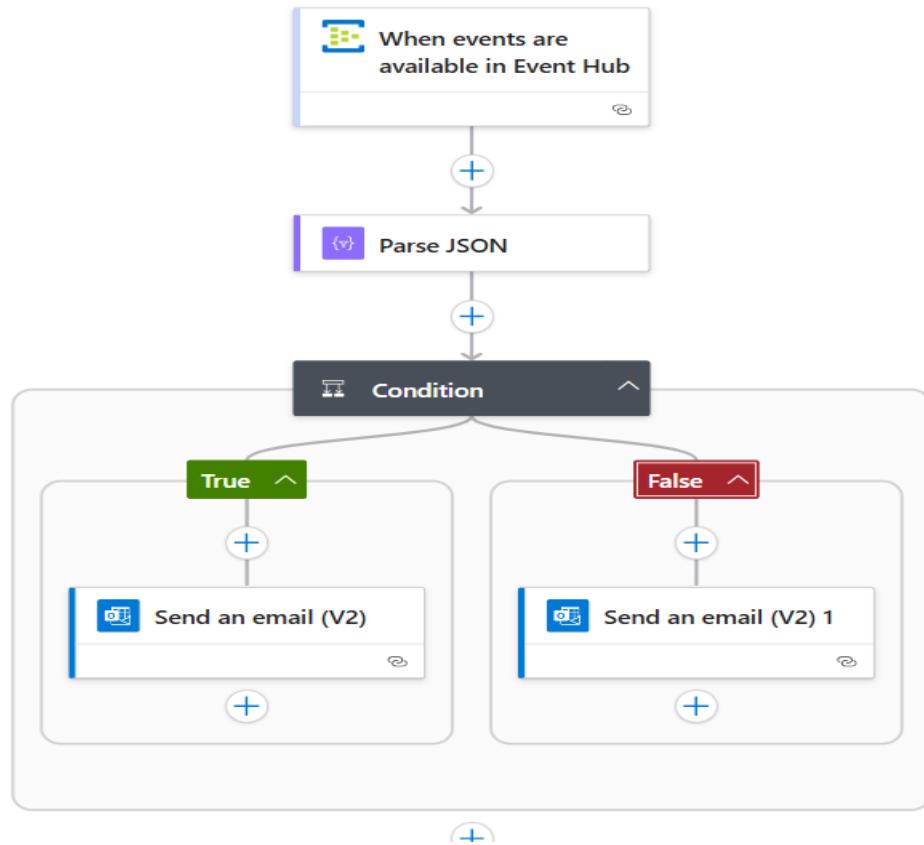
## 4.4 Logic App for Alerts

Triggers every 1 minute on alerts EventHub.

### Key Features:

- Parse JSON schema for event fields.
- Condition: If audience == 'citizen', send user-friendly email with Maps link; else, gov report with recommendations.
- High-importance HTML emails via Office 365.

## Figure 4.3: Logic App Flow Diagram



## 4.5 Genie AI Agent

Prompt-based on Gold layer for NLQ.

### General Instructions Snippet

#### Figure 4.4: Genie Query Example

What tables are there and how are they connected? Give me a short summary.

There are five main tables: fact\_traffic\_air (traffic and air quality sensor data), dim\_location (location details), dim\_sensor (sensor metadata), dim\_time (timestamp breakdown), and kpi\_city\_hour (hourly city-level KPIs). The tables are connected via keys: fact\_traffic\_air links to dim\_location and dim\_sensor by location\_id and sensor\_id, and to dim\_time by timestamp. kpi\_city\_hour provides aggregated city/hour metrics, making it easy to analyze trends by city and time.

## 4.6 Power BI Dashboards

2 pages imported from DataBricks SQL endpoint.

- **Traffic Page:**

- **Air Page:**

---

## 5. Testing and Results

### 5.1 Performance Testing

- **Ingestion:** 12K events/hour; 0 deserialization errors (EventHub dashboard).
- **Processing:** 45 alerts/min, 21 batches/hour; watermark delay: 5s (alerts), 31min (batches).

- ETL: ~10K Bronze rows → 8K Silver (20% dedup) → 24-hour KPIs in Gold.

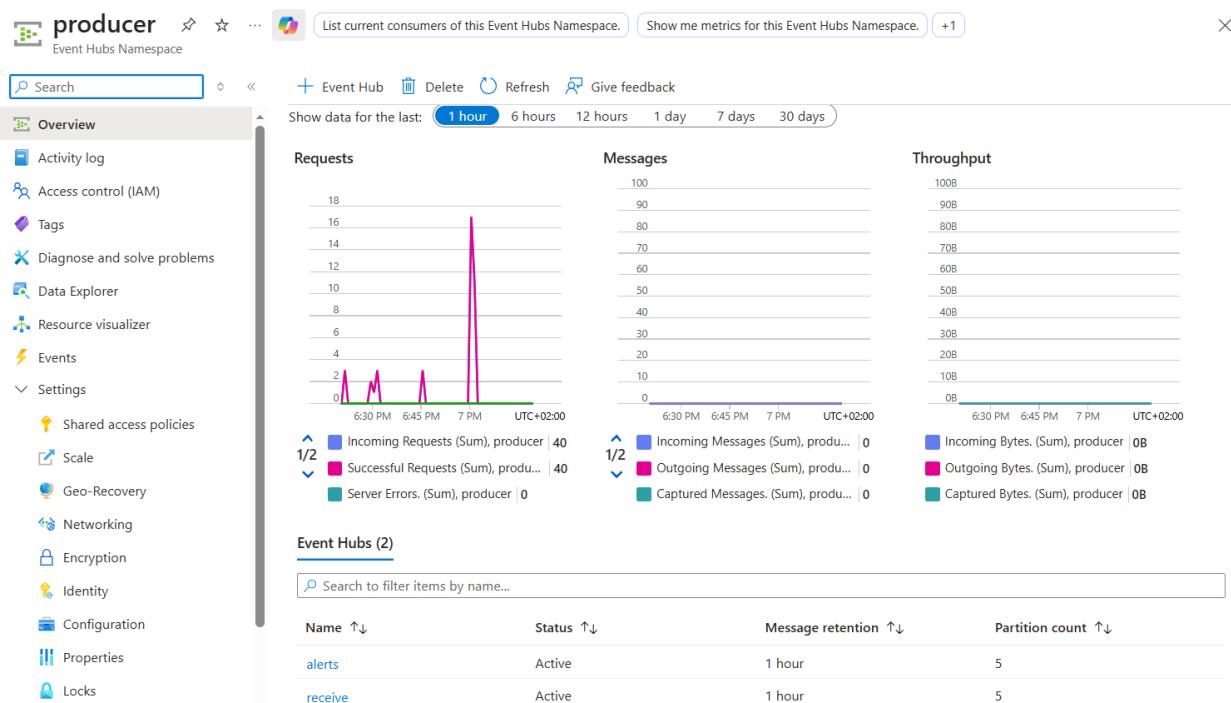
## 5.2 Key Insights

- Strong positive correlation (0.75) between congestion and NO<sub>2</sub>—validates traffic-emission hypothesis.
- Peak hours (7-10 AM, 4-7 PM) show 40% higher AQI; weekends reduce vehicles by 30%.
- Genie Example: Query "Traffic density in Giza?" → "Avg 80 vehicles/km during rush hour."

## 5.3 Challenges and Mitigations

- **Challenge:** High-volume streams → **Mitigation:** Partition by city in Stream Analytics.
- **Challenge:** Schema drift → **Mitigation:** Delta evolution in Silver.

**Figure 4.1: EventHub Metrics Over Time**



**Figure 4.2: Alert Email Sample**

## Government Alert Report: Unhealthy AQI – Cairo (Severity 5/5)



سعید علاء المتولى الغرباوي

سعید علاء المتولى الغرباوي

To: سعید علاء المتولى الغرباوي

😊 ⏪ ⏴ ⏵ ⏶ | 📎 ...

Sun 11/30/2025 1:07 PM

❗ High importance

🇬🇧 This message is in English

[Translate to Arabic](#) [Never translate from English](#)

**Dear Government Team,**

This is an automated **alert** from the Smart City Monitoring System.

An alert has been triggered in:

**City:** Cairo

**Sensor ID:** T-001

**Time:** 2025-11-30T13:04:19.638000Z

**Events in this window:**

**ALERT SUMMARY:**

**Alert Type:** Unhealthy AQI

**Alert Value:** 153

**Severity:** 5 / 5

**Status:** CRITICAL

**Message:**

Air quality is unhealthy!

Threshold Exceeded: Yes

Threshold Details:

AQI > 150

**RECOMMENDATIONS:**

Issuing a health advisory for affected districts

## Smart City Alert: High Accident Risk in Alex



سعید علاء المتولى الغرباوي

سعید علاء المتولى الغرباوي

To: سعید علاء المتولى الغرباوي

😊 ⏪ ⏴ ⏵ ⏶ | 📎 ...

Sun 11/30/2025 1:05 PM

❗ High importance

🇬🇧 This message is in English

[Translate to Arabic](#) [Never translate from English](#)

Hello ,

A real-time **alert** has been detected in your area in Alex.

**Alert Type:** High Accident Risk

**Current Value:** 0.69

**Severity:** 5 / 5

**Message:** High accident probability!

**Sensor ID:** T-003

**Time:** 2025-11-30T15:04:49.117921

Drive cautiously. Accident probability is high in your area.

Useful Link:

Check your live map **here**:

<https://maps.google.com/?q=Alex>

Stay safe,

Smart City Alert System

Reply

Forward

---

## 6. Recommendations and Future Work

### 6.1 Deployment Recommendations

- Use Terraform for IaC (see /deployment/terraform/main.tf).
- Monitor with Azure Sentinel (SLAs: <10s end-to-end).
- Scale: Add DLT for unified streaming; RBAC on Unity Catalog.

### 6.2 Enhancements

- **ML Layer:** LSTM on Gold for congestion forecasts (integrate Azure ML).
- **Outputs:** SMS alerts via Twilio; embed Power BI in citizen app.
- **Expansion:** Real sensors; multi-cloud (e.g., AWS Kinesis fallback).

### 6.3 Conclusion

This platform transforms raw streams into actionable intelligence, proving Azure's power for smart cities. It reduces risks and boosts efficiency ready for pilot in Cairo.